

Data Visualization BTech Computer Science Stream , January 2025 Week 5 - Data Cleaning and Preparation - Demonstration notebook Archana Praveek Kumar, Reg Number , Date: 06/01/2025

- During data analysis and modeling, a significant amount of time is spent on data preparation.
- Tasks include loading, cleaning, transforming, and rearranging data.
- Such tasks are often reported to take up 80% or more of an analyst's time.
- Pandas, along with the built-in Python language features, provides you with a high-level, flexible, and fast set of tools to enable you to manipulate data into the right form.

```
In [1]: import numpy as np
import pandas as pd
```

```
In [4]: # For float64 dtype, pandas uses NaN (Not a Number) to represent missing data
# Call this a sentinel value: when present, it indicates a missing (or null) value:
float_data = pd.Series([1.2, -3.5, np.nan, 0]) # to represent NA (Not available data)
float_data
#float_data.isna() # to check whether any value is NA or not
#float_data.notna()
```

```
Out[4]: 0      True
1      True
2     False
3      True
dtype: bool
```

```
In [3]: # The built-in Python None value is also treated as NA:
string_data = pd.Series(["aardvark", np.nan, None, "avocado"])
string_data
#string_data.isna()
```

```
Out[3]: 0    aardvark
1      NaN
2      None
3    avocado
dtype: object
```

- NA data may either be data that does not exist or that exists but was not observed.
- When cleaning up data for analysis, it is important to analyse missing data
- to identify data collection problems or potential biases in the data caused by missing data.

```
In [2]: # Strategies with missing data include 1. Dropping missing values
data1 = pd.Series([1, np.nan, 3.5, np.nan, 7])
data1
# by default drops any row containing a missing value
data1.dropna()
```

```
Out[2]: 0    1.0
2    3.5
4    7.0
dtype: float64
```

```
In [9]: data2 = pd.DataFrame([[1., 6.5, 3.], [1., np.nan, np.nan],
                           [np.nan, np.nan, np.nan], [np.nan, 6.5, 3.]])
data2
#data2.dropna(how="all") # drops rows that are all NA
data2.dropna(axis="columns", how="all")# filter to removes columns which have all valu
```

Out[9]:

	0	1	2
<b>0</b>	1.0	6.5	3.0
<b>1</b>	1.0	NaN	NaN
<b>2</b>	NaN	NaN	NaN
<b>3</b>	NaN	6.5	3.0

```
In [15]: # Strategies with missing data include 2. Filling missing values
df3 = pd.DataFrame(np.random.standard_normal((7, 3)))
df3.iloc[:4, 1] = np.nan
df3.iloc[:2, 2] = np.nan
df3.fillna(0)
#df3
```

Out[15]:

	0	1	2
<b>0</b>	-0.409814	0.000000	0.000000
<b>1</b>	1.278383	0.000000	0.000000
<b>2</b>	-1.377074	0.000000	-0.073055
<b>3</b>	0.217108	0.000000	0.912957
<b>4</b>	0.301200	-0.475921	-0.078225
<b>5</b>	0.958026	-0.829628	0.194278
<b>6</b>	-1.102895	-1.917630	0.388435

```
In [17]: # A dictionary where keys are column labels, and values are the fill values for those
df4 = pd.DataFrame(np.random.standard_normal((7, 3)))
df4.iloc[:4, 1] = np.nan
df4.iloc[:2, 2] = np.nan
#df4.fillna({1: 0.5, 2: 0})
#df4.fillna(df4.mean())
```

Out[17]:

	0	1	2
<b>0</b>	-0.105383	-0.266055	-0.106236
<b>1</b>	-0.799486	-0.266055	-0.106236
<b>2</b>	-0.397134	-0.266055	-0.784092
<b>3</b>	-0.392168	-0.266055	-1.299590
<b>4</b>	-0.356177	-1.420842	0.678405
<b>5</b>	-0.425674	1.203101	-1.024148
<b>6</b>	-0.068057	-0.580424	1.898248

In [7]:

```
# Can also use the replace method to replace 1 or more values.
data = pd.DataFrame({
    'A': [1, -999, 3, -1000],
    'B': [-1000, 5, -999, 7]
})
print(data)
# valid way to replace multiple values (-999 and -1000) with a single value (np.nan),
# If you want to modify the DataFrame or Series in place, use this argument. Otherwise,
data.replace([-999.0,-1000], np.nan, inplace=True)
print(data)
```

	A	B
0	1	-1000
1	-999	5
2	3	-999
3	-1000	7

  

	A	B
0	1.0	NaN
1	NaN	5.0
2	3.0	NaN
3	NaN	7.0

In [32]:

```
# Data Transformations include 1. removing duplicates
data5 = pd.DataFrame({"k1": ["one", "two"] * 3 + ["two"],
                      "k2": [1, 1, 2, 3, 3, 4, 4]})

#data5.duplicated()
data5.drop_duplicates()
#data5.drop_duplicates(subset=["k1"])
```

Out[32]:

	k1	k2
<b>0</b>	one	1
<b>1</b>	two	1
<b>2</b>	one	2
<b>3</b>	two	3
<b>4</b>	one	3
<b>5</b>	two	4
<b>6</b>	two	4

```
In [2]: # Data Transformations include 2. creating new columns
# DataFrame named data with two columns: "food" and "ounces"

data = pd.DataFrame({"food": ["salt", "bread", "butter",
                             "rice", "chips", "biscuits",
                             "chocolate", "honey", "sugar"],
                     "ounces": [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})

data
```

Out[2]:

	food	ounces
0	salt	4.0
1	bread	3.0
2	butter	12.0
3	rice	6.0
4	chips	7.5
5	biscuits	8.0
6	chocolate	3.0
7	honey	5.0
8	sugar	6.0

```
In [5]: # The dictionary food_to_type maps different types of food to their food type.

food_to_type = {
    "salt": "condiment",
    "bread": "breakfast",
    "butter": "condiment",
    "rice": "grocery",
    "chips": "snack",
    "honey": "condiment",
    "sugar": "condiment",
    "chocolate": "snack"
}
```

```
In [6]: #used to create a new column "ftype" in the data DataFrame, based on the values in the
#The map() function is applied to the "food" column, using the food_to_type dictionary

data["ftype"] = data["food"].map(food_to_type)
data
```

	food	ounces	ftype
0	salt	4.0	condiment
1	bread	3.0	breakfast
2	butter	12.0	condiment
3	rice	6.0	grocery
4	chips	7.5	snack
5	biscuits	8.0	NaN
6	chocolate	3.0	snack
7	honey	5.0	condiment
8	sugar	6.0	condiment

```
In [8]: # Data Transformations include 3. Discretization and Binning
# Discretization and Binning
ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
```

```
In [9]: bins = [18, 25, 35, 60, 100] #Assign custom Labels to each bin:
age_categories = pd.cut(ages, bins)#Assigns each age to one of these intervals. If an
age_categories
```

```
Out[9]: [(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100], (35, 6
0], (35, 60], (25, 35]]
Length: 12
Categories (4, interval[int64, right]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

```
In [10]: age_categories.codes# Returns the integer codes representing the bin each age belongs
print(age_categories.categories)#Returns the IntervalIndex of the categories (the bins
print(age_categories.categories[0])#Accesses the first category (interval
print(pd.value_counts(age_categories))

IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]], dtype='interval[int64, righ
t]')
(18, 25]
(18, 25]      5
(25, 35]      3
(35, 60]      3
(60, 100]     1
dtype: int64
```

```
In [11]: #By default, intervals are right-closed (e.g., (18, 25]), meaning the right endpoint i
#the intervals become left-closed (e.g., [18, 25)), meaning the left endpoint is inclu
pd.cut(ages, bins, right=False)
```

```
Out[11]: [[18, 25), [18, 25), [25, 35), [25, 35), [18, 25), ..., [25, 35), [60, 100), [35, 6
0), [35, 60), [25, 35]]
Length: 12
Categories (4, interval[int64, left]): [[18, 25) < [25, 35) < [35, 60) < [60, 100]]
```

```
In [12]: #Assigns custom Labels to the bins:
#(18, 25] → "Youth"
#(25, 35] → "YoungAdult"
#(35, 60] → "MiddleAged"
#(60, 100] → "Senior"
```

```
#Each age is categorized into one of the groups based on the bin it falls into.
#for example:
#20 → "Youth"
#27 → "YoungAdult"
#61 → "Senior"
```

In [13]: `group_names = ["Youth", "YoungAdult", "MiddleAged", "Senior"]  
pd.cut(ages, bins, labels=group_names)`

Out[13]: `['Youth', 'Youth', 'Youth', 'YoungAdult', 'Youth', ..., 'YoungAdult', 'Senior', 'MiddleAged', 'MiddleAged', 'YoungAdult']  
Length: 12  
Categories (4, object): ['Youth' < 'YoungAdult' < 'MiddleAged' < 'Senior']`

In [14]: `print(pd.value_counts(age_categories))`

```
(18, 25]      5
(25, 35]      3
(35, 60]      3
(60, 100]     1
dtype: int64
```

In [49]: `# Data Transformations includes 4. Detecting and Filtering Outliers:  
data = pd.DataFrame(np.random.standard_normal((1000, 4)))  
data.describe()`

Out[49]:

	<b>0</b>	<b>1</b>	<b>2</b>	<b>3</b>
<b>count</b>	1000.000000	1000.000000	1000.000000	1000.000000
<b>mean</b>	0.035320	-0.047482	-0.040091	0.006336
<b>std</b>	0.969095	0.962987	0.994311	1.019735
<b>min</b>	-3.333767	-3.194414	-3.108915	-3.645860
<b>25%</b>	-0.619875	-0.735556	-0.734509	-0.680576
<b>50%</b>	0.034764	-0.048538	-0.049977	0.040150
<b>75%</b>	0.688640	0.647824	0.621226	0.713369
<b>max</b>	3.525865	2.611678	3.366626	2.763474

In [51]: `col = data[2] #Extracts the third column (index 2) of the DataFrame data, Computes the col[col.abs() > 3] #Filters the col Series, keeping only the rows where the absolute v`

Out[51]: `134 3.366626
630 -3.108915
Name: 2, dtype: float64`

In [53]: `#data.abs(): Computes the absolute value of all elements in the DataFrame.  
# returns a DataFrame of the same shape with True for values where the condition is met  
.any(axis="columns"): Checks each row across all columns to see if any value in that data[(data.abs() > 3).any(axis="columns")]`

Out[53]:

	0	1	2	3
<b>42</b>	0.208011	-0.150923	-0.362528	-3.548824
<b>134</b>	0.193299	1.397822	3.366626	-2.372214
<b>281</b>	3.525865	0.283070	0.544635	0.462204
<b>301</b>	-0.450721	-0.080332	0.599947	-3.645860
<b>524</b>	-3.333767	-1.240685	-0.650855	0.076254
<b>605</b>	0.344072	0.581893	-1.116332	-3.018842
<b>630</b>	-0.555434	-0.048478	-3.108915	1.117755
<b>760</b>	-0.217146	-0.274138	1.188742	-3.183867
<b>807</b>	0.744019	1.741426	-2.214074	-3.140963
<b>973</b>	-0.848098	-3.194414	0.077839	-1.733549

In [54]:

```
data[data.abs() > 3] = np.sign(data) * 3 # Cap values at ±3
#np.sign(data) * 3: Multiplies the sign of each value by 3: Values greater than 3 are
data.describe()
```

Out[54]:

	0	1	2	3
<b>count</b>	1000.000000	1000.000000	1000.000000	1000.000000
<b>mean</b>	0.035128	-0.047288	-0.040348	0.007874
<b>std</b>	0.966234	0.962371	0.992790	1.014803
<b>min</b>	-3.000000	-3.000000	-3.000000	-3.000000
<b>25%</b>	-0.619875	-0.735556	-0.734509	-0.680576
<b>50%</b>	0.034764	-0.048538	-0.049977	0.040150
<b>75%</b>	0.688640	0.647824	0.621226	0.713369
<b>max</b>	3.000000	2.611678	3.000000	2.763474

In [63]:

```
# Data Transformation include 5. Computing indicator/dummy variables.
#key: Categorical column with values a, b, and c.
#data1: Numerical column with values from 0 to 5.
df = pd.DataFrame({"key": ["b", "b", "a", "c", "a", "b"],
                    "data1": range(6)})
print(df)
pd.get_dummies(df["key"], dtype=float)
```

	key	data1
0	b	0
1	b	1
2	a	2
3	c	3
4	a	4
5	b	5

Out[63]:

	<b>a</b>	<b>b</b>	<b>c</b>
<b>0</b>	0.0	1.0	0.0
<b>1</b>	0.0	1.0	0.0
<b>2</b>	1.0	0.0	0.0
<b>3</b>	0.0	0.0	1.0
<b>4</b>	1.0	0.0	0.0
<b>5</b>	0.0	1.0	0.0